

# AlexNet



ImageNet Classification  
with Deep Convolutional Neural Networks

Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton



Lisa Fan & Jason Krone

# Introduction

- Submitted to ImageNet Challenge in 2012
  - Won competition with 16% top 5 error. Large improvement over 2nd runner-up with 26%
- Popularized CNNs for computer vision
  - Although CNNs were previously used, e.g. LeNet (Yann LeCun), AlexNet was deeper, bigger, and included stacked Convolutional layers.
- Introduced ReLU activation function

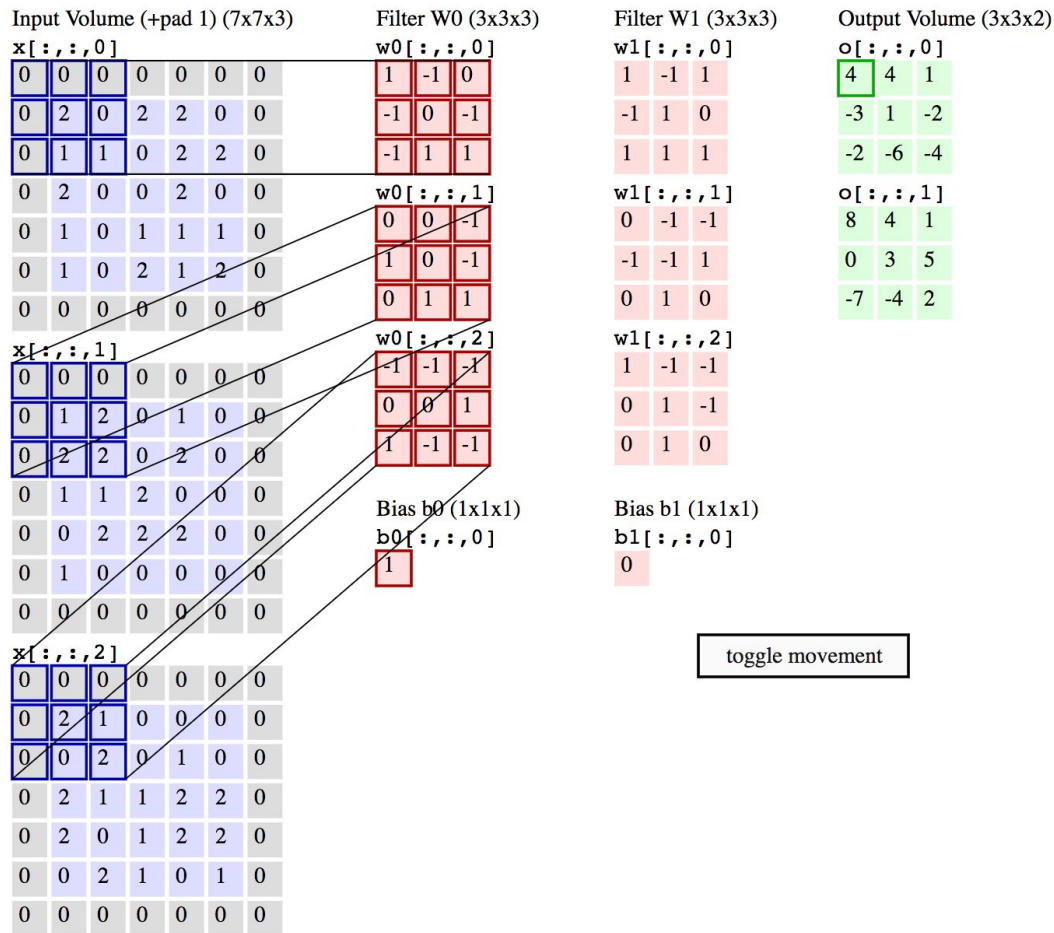


# Setting future trends: GPUs + Deep Learning




# Aside: Convolution Layers

- Element-wise multiplication
- Add all terms
- $x * w[:, :, 0] =$
- $x * w[:, :, 1] =$
- $x * w[:, :, 2] =$
- $x * w + \text{bias} =$



# Architecture

- 8 layers: 5 Conv layers, 3 Fully connected (FC)
- Output is fed to a 1000-way softmax function




# Training on GPUs

Intra GPU connection:



2nd, 4th, and 5th layers

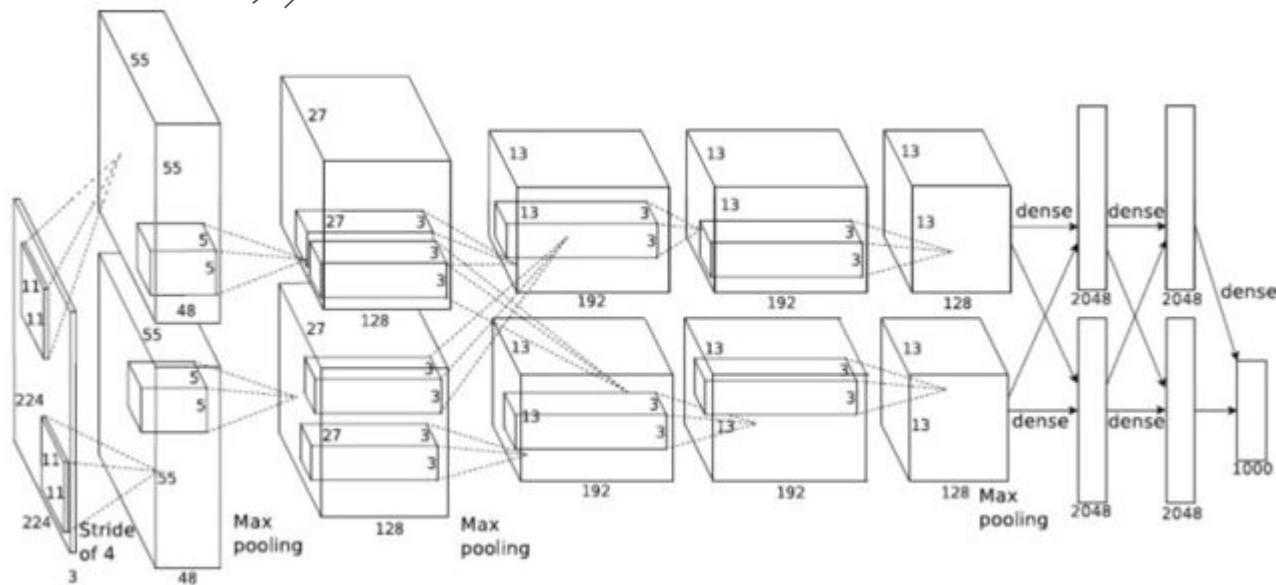
Inter GPU connection:





3rd, 6th, 7th, 8th layers

GPU #1

GPU #2





# Performance: CPUs vs GPUs



# First Convolution Layer

- Input: 224x224x3
- 96 11x11x3 filters with stride of 4
- Output?  $(W - F) / S + 1$






# Overlapping Pooling

- AlexNet uses **overlapping max pooling** after the first and second convolution layer
- The “pooling filter” is 3x3 with a stride of 2
- Input: 55x55x48
- Output?  $(W - F + 2P) / S + 1$

1	2	2	3	4	5	8
7	6	1	2	0	7	9
2	3	4	5	6	1	2




7		
---	--	--





# Second Convolution Layer

- Layer 2 Input: 27x27x48 (on each GPU)
- 256 5x5x48 filters
- Padding = 2, stride = 1
- Output?  $(W - F + 2P) / S + 1$



# Third to Fifth Convolution Layers

- Stride = 1, Padding = 1
- Layer 3:
  - Input: 13x13x128 (on each GPU)
  - 384 3x3x256 filters
  - Output: 13x13x384 (or 13x13x192 on each GPU)
  - Cross-GPU Layer: All output goes through all filters
- Layer 4:
  - 384 3x3x192 filters
  - Output: 13x13x384 (or 13x13x192 on each GPU)
- Layer 5:
  - 256 3x3x192 filters
  - Output: 13x13x256 (or 13x13x128 on each GPU)
  - After pooling: 6x6x256



# Layers 6, 7, 8: Fully Connected Layers

- All FC Layers are cross-GPU: all inputs go through all neurons

- Layer 6


- Input: 6x6x256
- 4096 neurons
- Output: 4096

- Layer 7

- 4096 neurons


- Layer 8

- 1000 neurons
- Softmax to get final scores
- Final output: 1000x1 vector
- ImageNet has 1000 classes



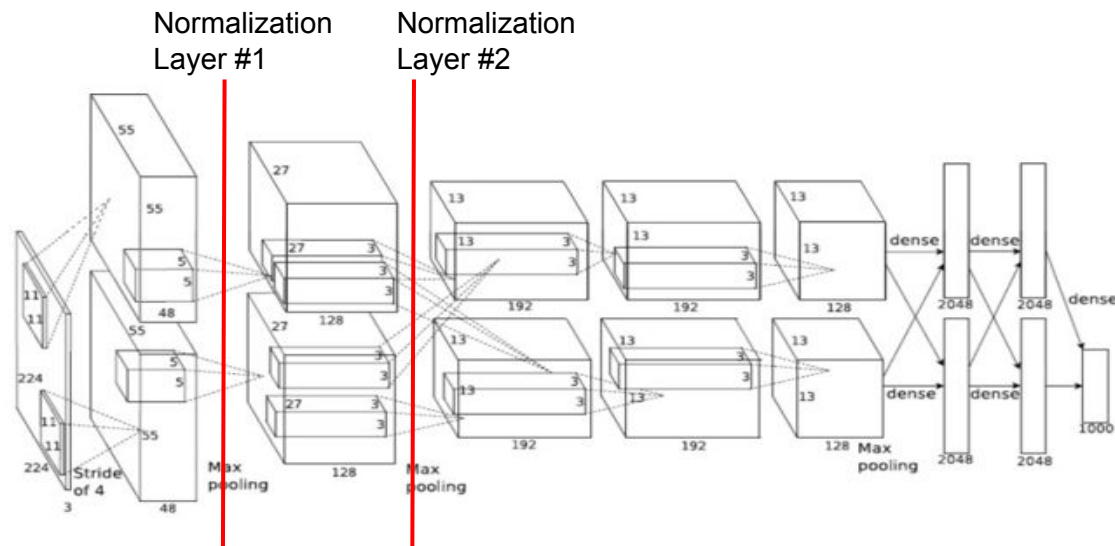
# Parameters in AlexNet: First Conv Layer

- 96  $11 \times 11 \times 3$  filters
- Output:  $55 \times 55 \times 96$
- How many parameters?
  - $11 \times 11 \times 3 = 363$  weights per filter
  - $55 \times 55 \times 96 = 290,400$  neurons
  - $290,400 \times 364 = 105,705,600$  parameters!
- Parameter sharing
  - Neurons of each depth slice share the same weights and bias
  - We now have  $96 \times 11 \times 11 \times 3 = 34,848$  parameters instead of ~100 million
- AlexNet ends up having 60 million parameters and 650,000 neurons!



# Response normalization

- Normalize after ReLU before pooling
- “Brightness normalization”
- Difference from batch normalization
  - Internal covariate shift



# Response normalization

b - response normalized activity

a - activity of a neuron by applying kernel (or filter) i at position (x, y)


N - number of kernels in a given layer

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^{\beta}$$

Hyper-parameters: k = 2, alpha = 10<sup>-4</sup>, beta = 0.75, n = 5

# Reducing Overfitting: Data Augmentation

- Take 5 224x224 patches from original 256x256 images
  - Each corner + center
  - Take horizontal reflections of each of these
  - Dataset is now 10x bigger
  - Average scores of 10 patches at test time
- PCA performed on RGB pixel values







# More Details

- Update rule: Stochastic Gradient Descent + Momentum
- Batch size: 128
- Weight Decay: 0.0005
- Learning rate: 0.01
- Weights initialized from Gaussian distribution (mean=0, std dev=0.01)
- Neuron biases initialized to 1 to provide ReLU with positive inputs
- Training set size: 1.2 million images
- Run for ~90 cycles
- Took 5~6 days to run

# Weaknesses



alexnet flaw



**flawless**

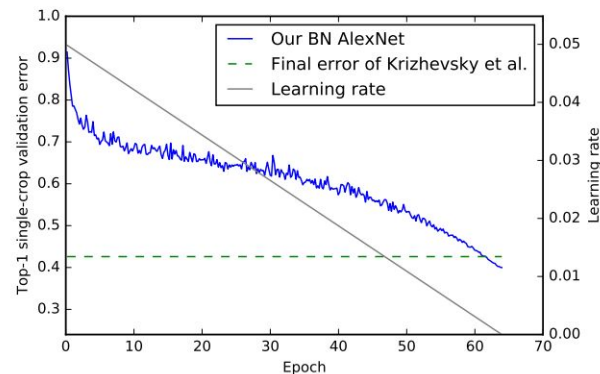
Press Enter to search.

# Room for Improvement

- Batch Normalization
  - Simon et. al. added a batch normalization layer between each convolutional and activation unit layer, and removed the local response normalization and dropout layers.
- Not great for object detection (bounding box)
- Shallow network by today's standards

Table 1. **Single-crop** top-1 and top-5 error of our models on the validation set of ILSVRC 2012.

Model	Top-1 error		Top-5 error	
	Ours	Original	Ours	Original
AlexNet	<b>39.9%</b>	42.6%	<b>18.1%</b>	19.6%
VGG19	<b>26.9%</b>	28.7%	<b>8.8%</b>	9.9%
ResNet-10	<b>36.1%</b>	—	<b>14.8%</b>	—
ResNet-50	<b>24.6%</b>	24.7%	<b>7.6%</b>	7.8%



# Questions

- How does each filter in a layer learn different weights?
  - Random Initialization



# Works Cited

- <http://cs231n.github.io/convolutional-networks/>
- [http://vision.stanford.edu/teaching/cs231b\\_spring1415/slides/alexnet\\_tugce\\_kyunghee.pdf](http://vision.stanford.edu/teaching/cs231b_spring1415/slides/alexnet_tugce_kyunghee.pdf)
- <https://www.analyticsvidhya.com/blog/2016/04/deep-learning-computer-vision-introduction-convolution-neural-networks/>
- Simon, Marcel, Erik Rodner, and Joachim Denzler. "ImageNet pre-trained models with batch normalization." arXiv preprint arXiv:1612.01452 (2016).
- Demo: [http://www.cs.toronto.edu/~guerzhoy/tf\\_alexnet/](http://www.cs.toronto.edu/~guerzhoy/tf_alexnet/)