

Lecture 12:

Activity Recognition and Unsupervised Learning

Tuesday April 4, 2017

Announcements!

- International Max Planck Research School for Intelligent Systems with director Michael Black, applications open for 100 new PhD students
- Final Project milestones due today
- Vote for Final Day and Location on Doodle, if you didn't get a Doodle link let me know
- Complain about AWS availability to t-staff

Activity Recognition

Classic Video Segmentation: **Optical Flow**

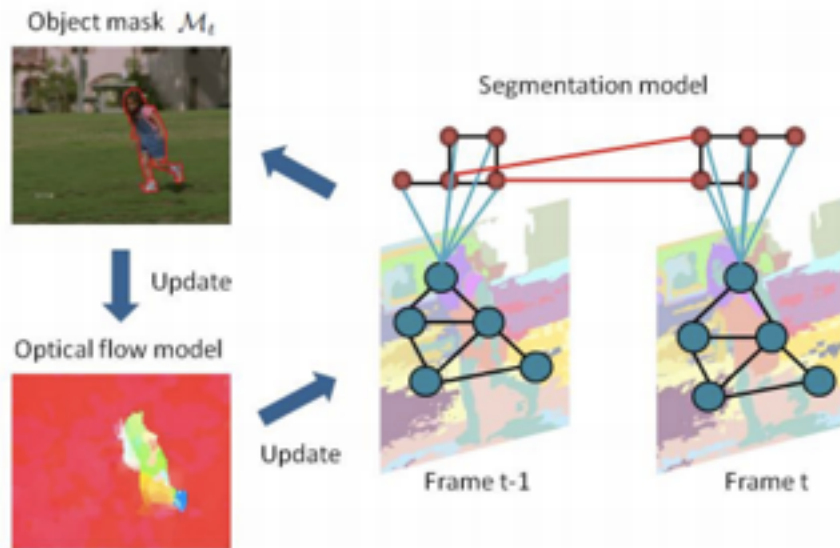



Figure 2. Overview of the proposed model. For segmentation, we consider a multi-level spatial-temporal model. Red circles denote pixels, which belong to the superpixel marked by the turquoise circles. The black and the red lines denote the spatial and temporal relationships, respectively. The relationships between the pixels and the superpixel are denoted by the turquoise lines. After obtaining the object mask, \mathcal{M}_t , we use this mask to re-estimate the optical flow, and update both models iteratively.

Latest Iteration:
Video Segmentation via object flow
Tsai et al., 2016

[G. Farnebäck, "Two-frame motion estimation based on polynomial expansion," 2003]
[T. Brox and J. Malik, "Large displacement optical flow: Descriptor matching in variational motion estimation," 2011]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4


=>

Output volume **[55x55x96]**

Q: What if the input is now a small chunk of video? E.g. [227x227x3x15] ?

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4

=>

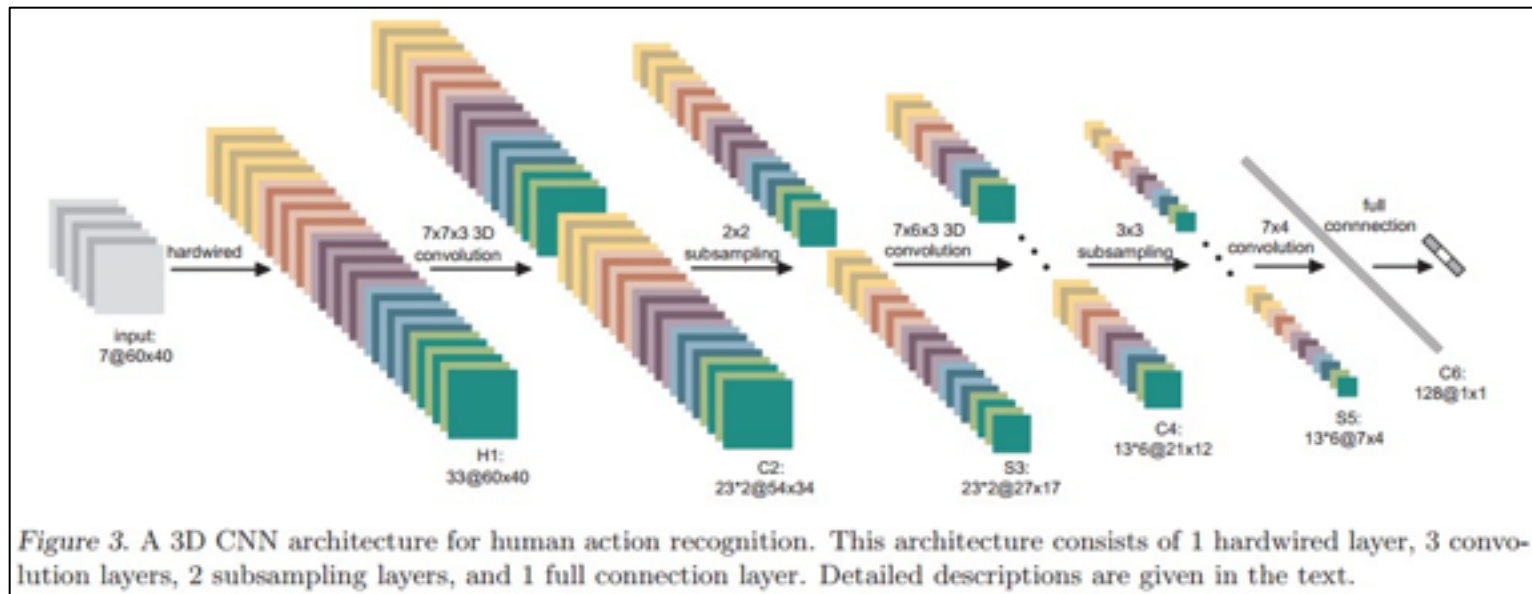
Output volume **[55x55x96]**

Q: What if the input is now a small chunk of video? E.g. [227x227x3x15] ?

A: Extend the convolutional filters in time, perform spatio-temporal convolutions!

E.g. can have 11x11xT filters, where T = 2..15.

Spatio-Temporal ConvNets



[3D Convolutional Neural Networks for Human Action Recognition, Ji et al., 2010]

Spatio-Temporal ConvNets


Learned filters on
the first layer



[Large-scale Video Classification with Convolutional Neural Networks, Karpathy et al., 2014]

Long-time Spatio-Temporal ConvNets

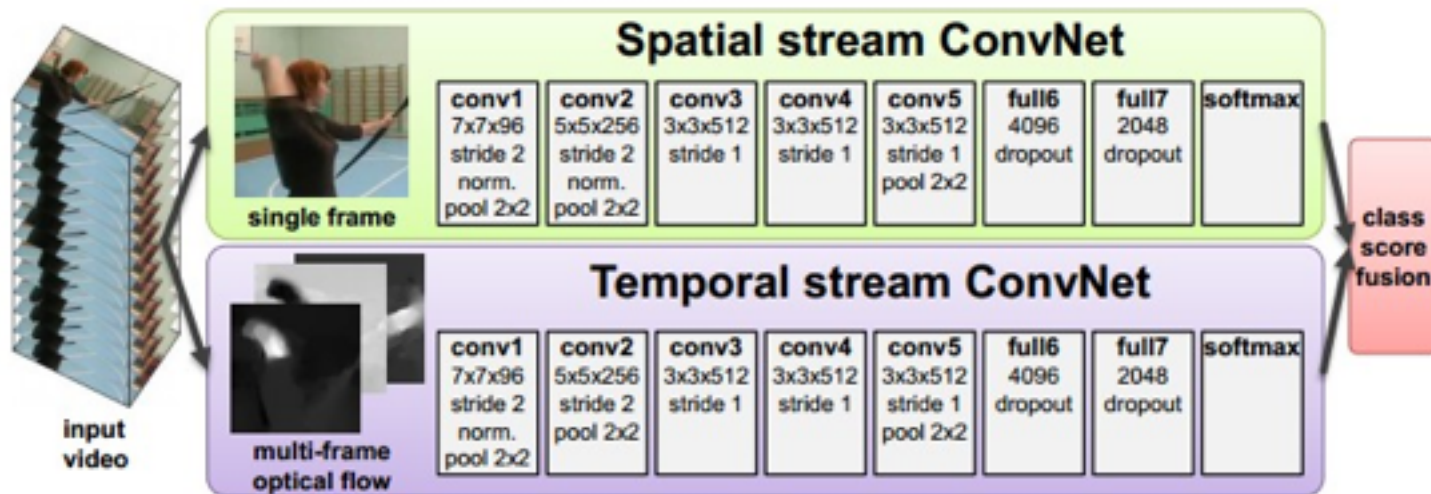
LSTM way before it was cool



(This paper was ahead of its time. Cited 65 times.)

Sequential Deep Learning for Human Action Recognition, Baccouche et al., 2011


Spatio-Temporal ConvNets



[Two-Stream Convolutional Networks for Action Recognition in Videos, Simonyan and Zisserman 2014]

[T. Brox and J. Malik, "Large displacement optical flow: Descriptor matching in variational motion estimation," 2011]

Spatio-Temporal ConvNets



Spatial stream ConvNet	73.0%	40.5%
Temporal stream ConvNet	83.7%	54.6%
Two-stream model (fusion by averaging)	86.9%	58.0%
Two-stream model (fusion by SVM)	88.0%	59.4%

Two-stream version works much better than either alone.


[Two-Stream Convolutional Networks for Action Recognition in Videos, **Simonyan** and Zisserman 2014]

[T. Brox and J. Malik, “Large displacement optical flow: Descriptor matching in variational motion estimation,” 2011]

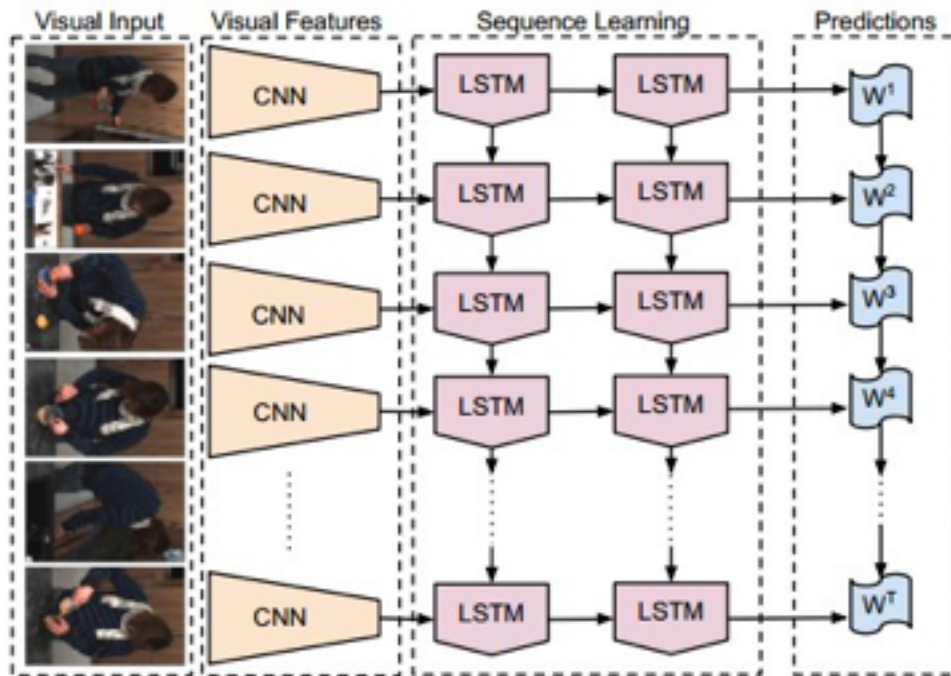
Long-time Spatio-Temporal ConvNets

All 3D ConvNets so far used local motion cues to get extra accuracy (e.g. half a second or so)

Q: what if the temporal dependencies of interest are much much longer? E.g. several seconds?

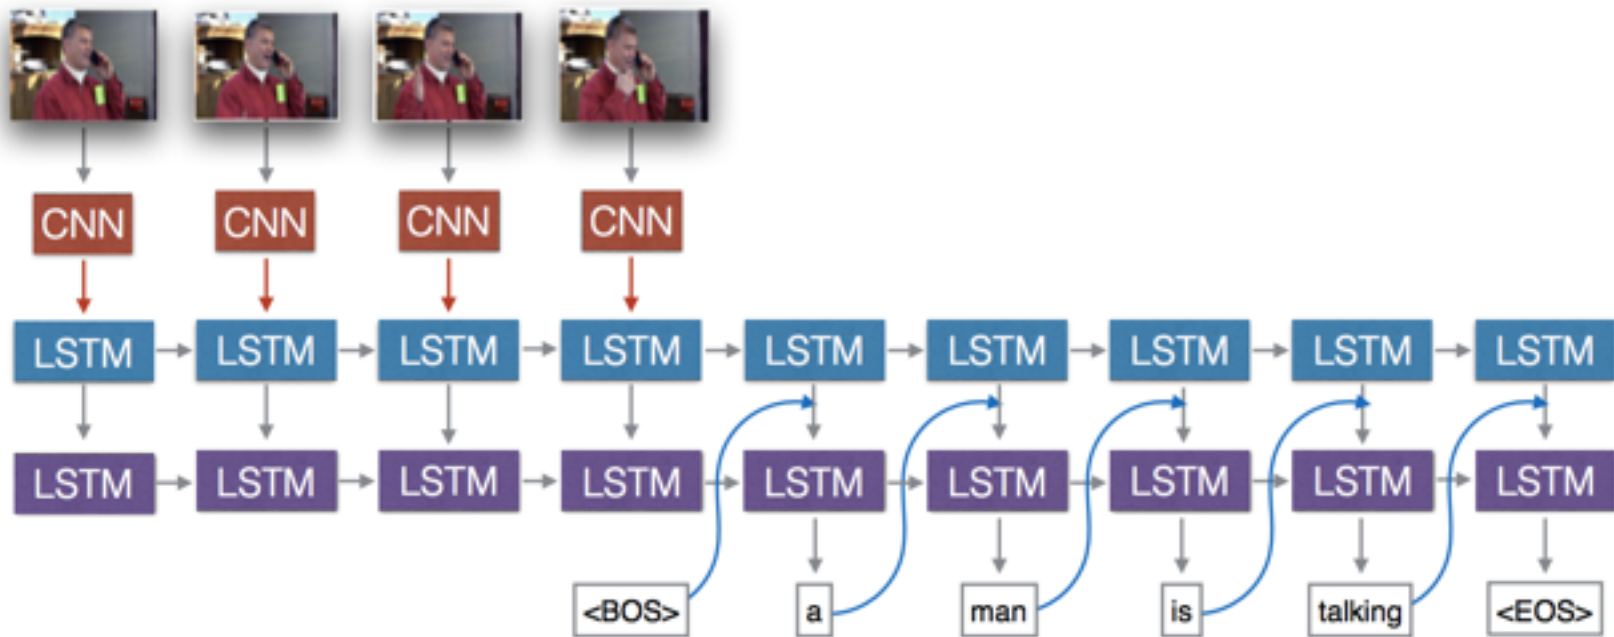


Long-time Spatio-Temporal ConvNets



[Long-term Recurrent Convolutional Networks for Visual Recognition and Description, Donahue et al., 2015]

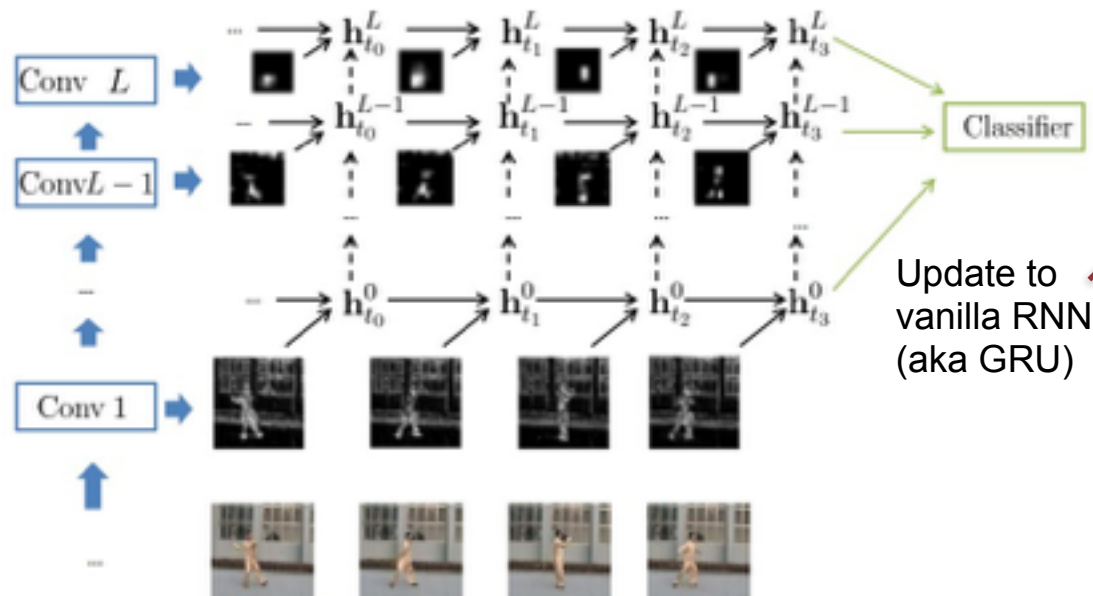
Video Description



Venugopalan et al., "Sequence to Sequence -- Video to Text," 2015.

sequential input & output

Long-time Spatio-Temporal ConvNets



All neurons in the ConvNet are recurrent.

$$\begin{aligned}
 z_t^l &= \sigma(W_z^l * x_t^l + U_z^l * h_{t-1}^l), & \text{update gate} \\
 r_t^l &= \sigma(W_r^l * x_t^l + U_r^l * h_{t-1}^l), & \text{reset gate} \\
 \tilde{h}_t^l &= \tanh(W^l * x_t^l + U * (r_t^l \odot h_{t-1}^l)), \\
 h_t^l &= (1 - z_t^l)h_{t-1}^l + z_t^l \tilde{h}_t^l,
 \end{aligned}$$

Update to
vanilla RNN
(aka GRU)

Only requires (existing) 2D
CONV routines. No need for 3D
spatio-temporal CONV.

[Delving Deeper into Convolutional Networks for Learning Video Representations, Ballas et al., 2016]

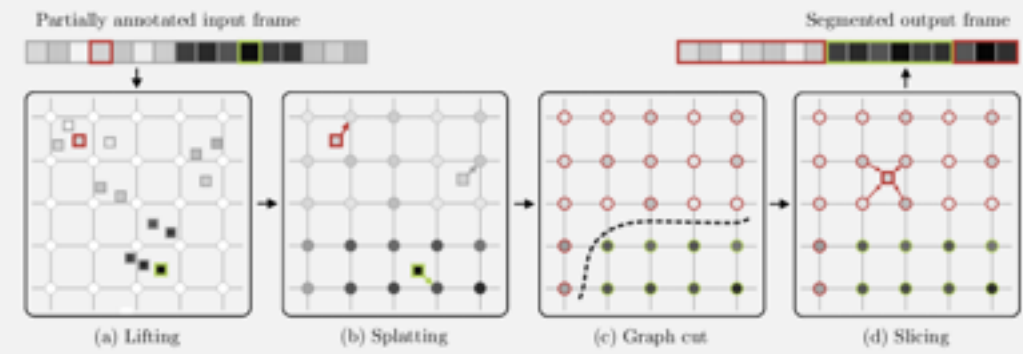
Propagation

Graph Cut for Video:

Bilateral Space Video Segmentation
Marki et al., 2016



The algorithm consists of 4 steps: lifting, splatting, graph cut, and slicing.



Unsupervised Learning

Unsupervised Learning Overview

- Autoencoders
 - Vanilla
 - Variational
- Adversarial Networks

Supervised vs Unsupervised

- **Supervised Learning**

- **Data:** (x, y)

- x is data, y is label

- **Goal:** Learn a *function* to map $x \rightarrow y$

- **Examples:** Classification, regression, object detection, semantic segmentation, image captioning, etc

Supervised vs Unsupervised

- Supervised Learning

- **Data:** (x, y)

- x is data, y is label

- **Goal:** Learn a *function* to map $x \rightarrow y$

- **Examples:** Classification, regression, object detection, semantic segmentation, image captioning, etc

Unsupervised Learning

Data: x

Just data, no labels!



Goal: Learn some *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, generative models, etc.

Unsupervised Learning

- Autoencoders
 - Traditional: feature learning
 - Variational: generate samples
- Generative Adversarial Networks: Generate samples

Autoencoders





Autoencoders

Originally: Linear + nonlinearity (sigmoid)

Later: Deep, fully-connected



Later: ReLU CNN






Autoencoders

\mathbf{z} usually smaller than \mathbf{x}
(dimensionality reduction)
Prevents trivial solution



Originally: Linear + nonlinearity (sigmoid)
Later: Deep, fully-connected
Later: ReLU CNN




Autoencoders



Autoencoders


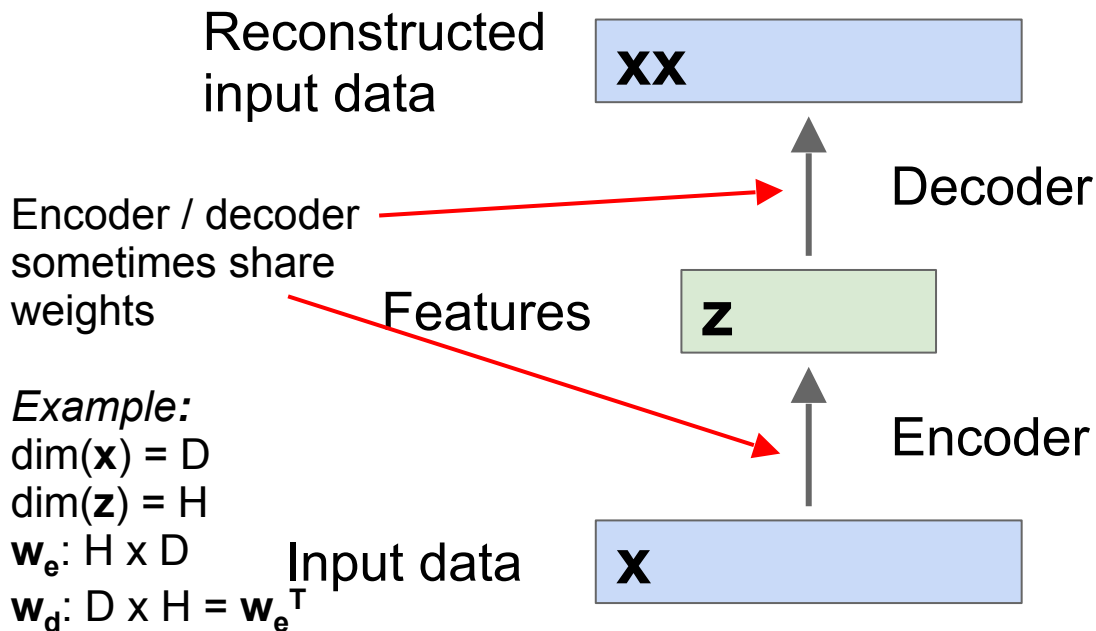


Encoder: 4-layer conv
Decoder: 4-layer upconv





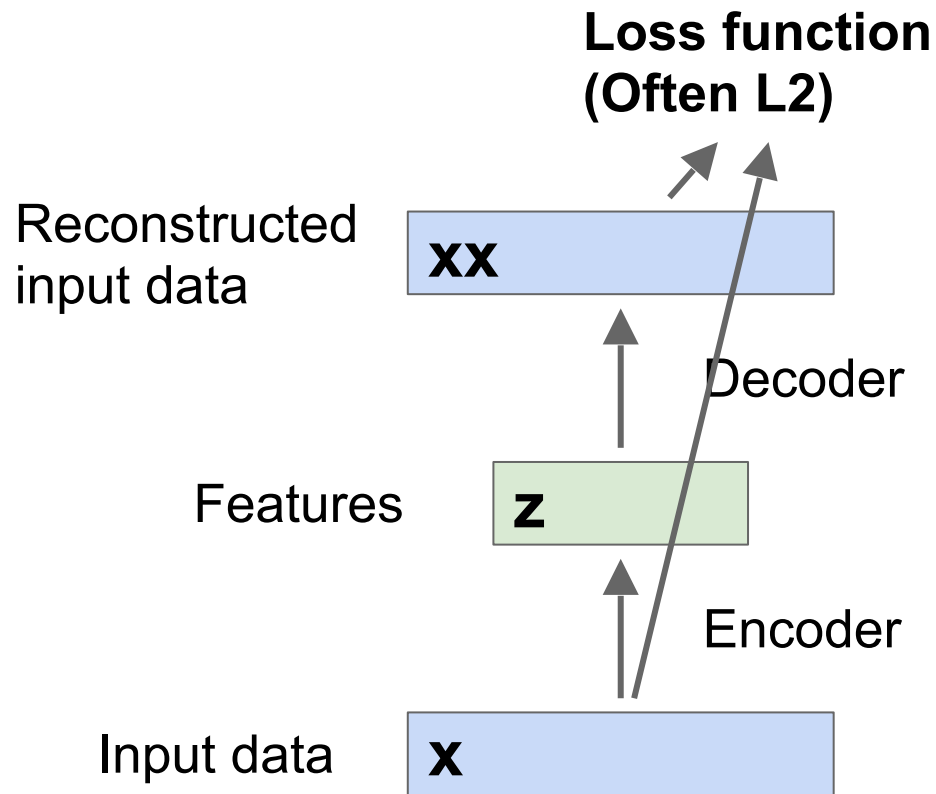
Autoencoders

Goal: Train for reconstruction with no labels!




Encoder: 4-layer conv
Decoder: 4-layer upconv



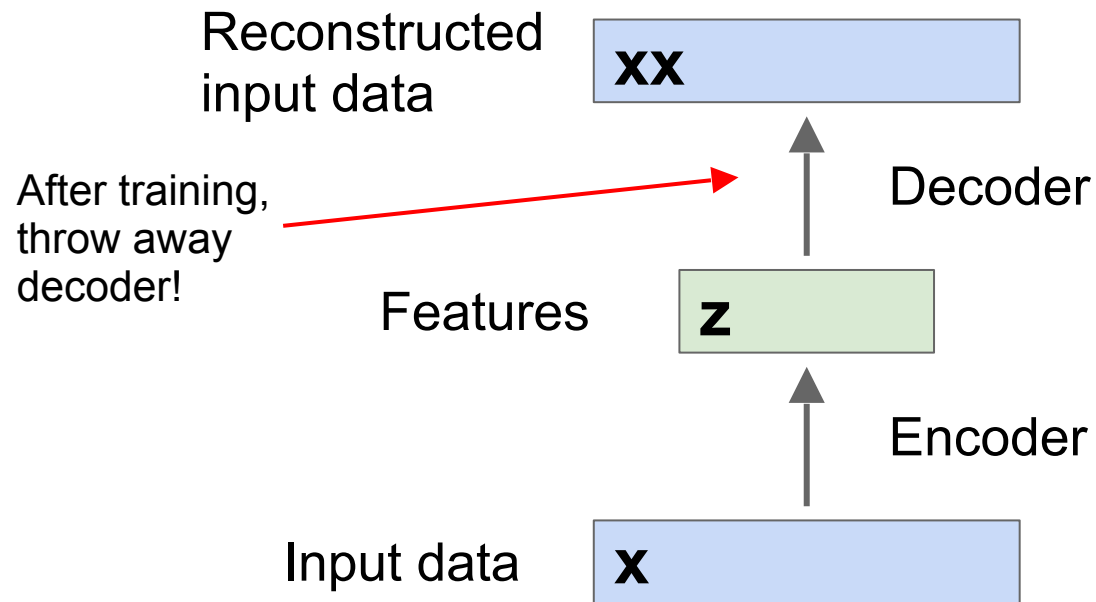


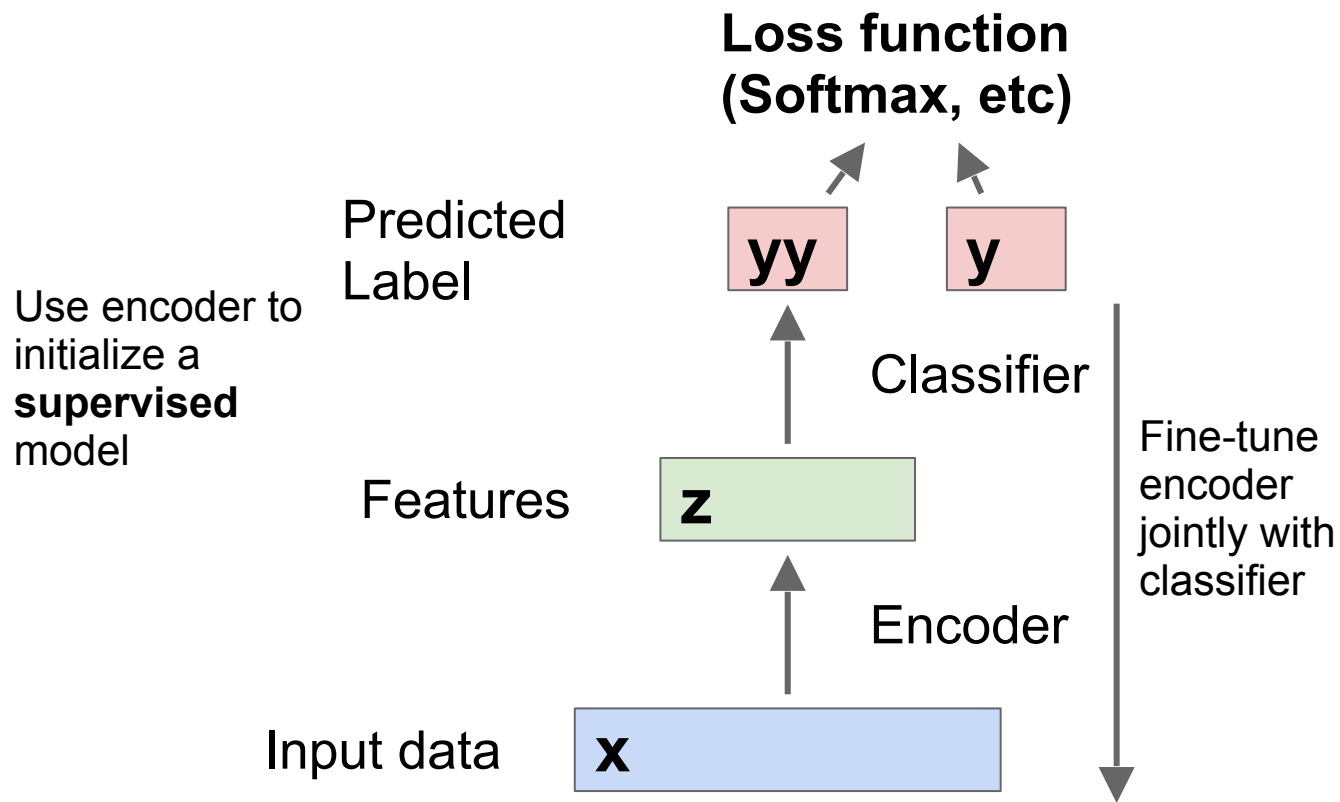


Train for
reconstruction
with no labels!




Autoencoders





bird plane
dog deer truck

Train for final task
(sometimes with
small data)

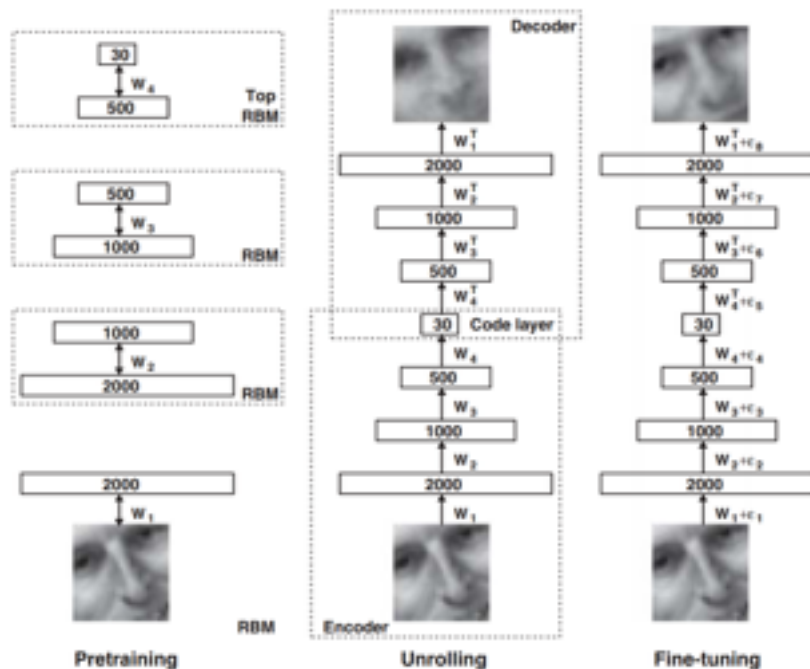


Autoencoders: Greedy Training

In mid 2000s layer-wise pretraining with Restricted Boltzmann Machines (RBM) was common

Training deep nets was hard in 2006!

It is difficult to optimize the weights in nonlinear autoencoders that have multiple hidden layers (2–4). With large initial weights, autoencoders typically find poor local minima; with small initial weights, the gradients in the early layers are tiny, making it infeasible to train autoencoders with many hidden layers. If



Not common anymore

With ReLU, proper initialization, batchnorm, Adam, etc easily train from scratch

Hinton and Salakhutdinov, "Reducing the Dimensionality of Data with Neural Networks", Science 2006

Alternatives




Fig. 1. Visualization of MNIST test set in a 2D space by classic 2-dimensional auto-encoder




Fig. 2. Visualization of MNIST test set in a 2D space by Siamese network

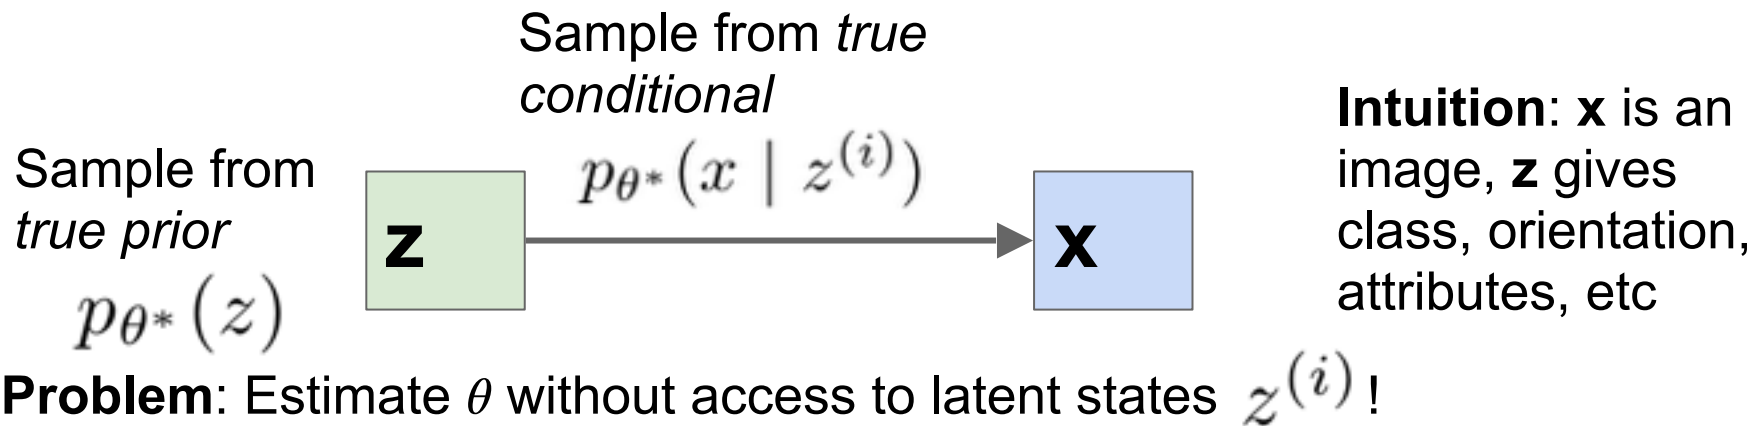
- Siamese Networks
- Triplet Networks
- Pretraining on unrelated supervised task (aka Transfer Learning)

Generating Samples

- What if you want to make new examples?
- Need Generative Model
- MCMC?
 - too slow, hard to scale
- MAP / Maximization?
 - Strong overfitting of high dimensional data — won't generate a large variety of interesting things

Variational Autoencoder a Generative Method

- A Bayesian spin on an autoencoder - lets us generate data!
- Assume our data $\{x^{(i)}\}_{i=1}^N$ is generated like this:



Variational Autoencoder: Encoder

- By Bayes Rule the posterior is:

$$p_{\theta}(z | x) = \frac{p_{\theta}(x | z) p_{\theta}(z)}{p_{\theta}(x)}$$


Use decoder network =)
Gaussian =)
Intractable integral =)

Approximate posterior with
encoder network $q_{\phi}(z | x)$

Fully-connected
or convolutional


Mean and (diagonal)
covariance of
 $q_{\phi}(z | x)$

Encoder network
with parameters ϕ



Variational auto-encoder

injected noise



$p(z)$ and $p(x|z)$
(decoder)

$q(z|x) = \mathcal{N}(\mu, \sigma)$
(encoder)

Solution: Approximate
posterior with **encoder
network** $q_{\phi}(z | x)$

Variational Autoencoder a Generative Method

Key reparameterization trick

Construct samples $z \sim q_\varphi(z|x)$ in two steps:

1. $\varepsilon \sim p(\varepsilon)$ (*random seed independent of φ*)
2. $z = g(\varphi, \varepsilon, x)$ (*differentiable perturbation*)

such that $z \sim q_\varphi(z|x)$ (*the correct distribution*)

Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014

Auto-Encoding Variational Bayes

Online algorithm

repeat

$\mathbf{x} \leftarrow$ random datapoint or minibatch

$\epsilon \leftarrow$ sample from $p(\epsilon)$

$g_{\theta}, g_{\phi} \leftarrow \nabla_{\theta, \phi} \tilde{\mathcal{L}}(\theta, \phi; \mathbf{x}, g(\epsilon, \phi))$

Decoder Network Parameters $\theta \leftarrow \theta + \alpha \cdot g_{\theta}$

Encoder Network Parameters $\phi \leftarrow \phi + \alpha \cdot g_{\phi}$

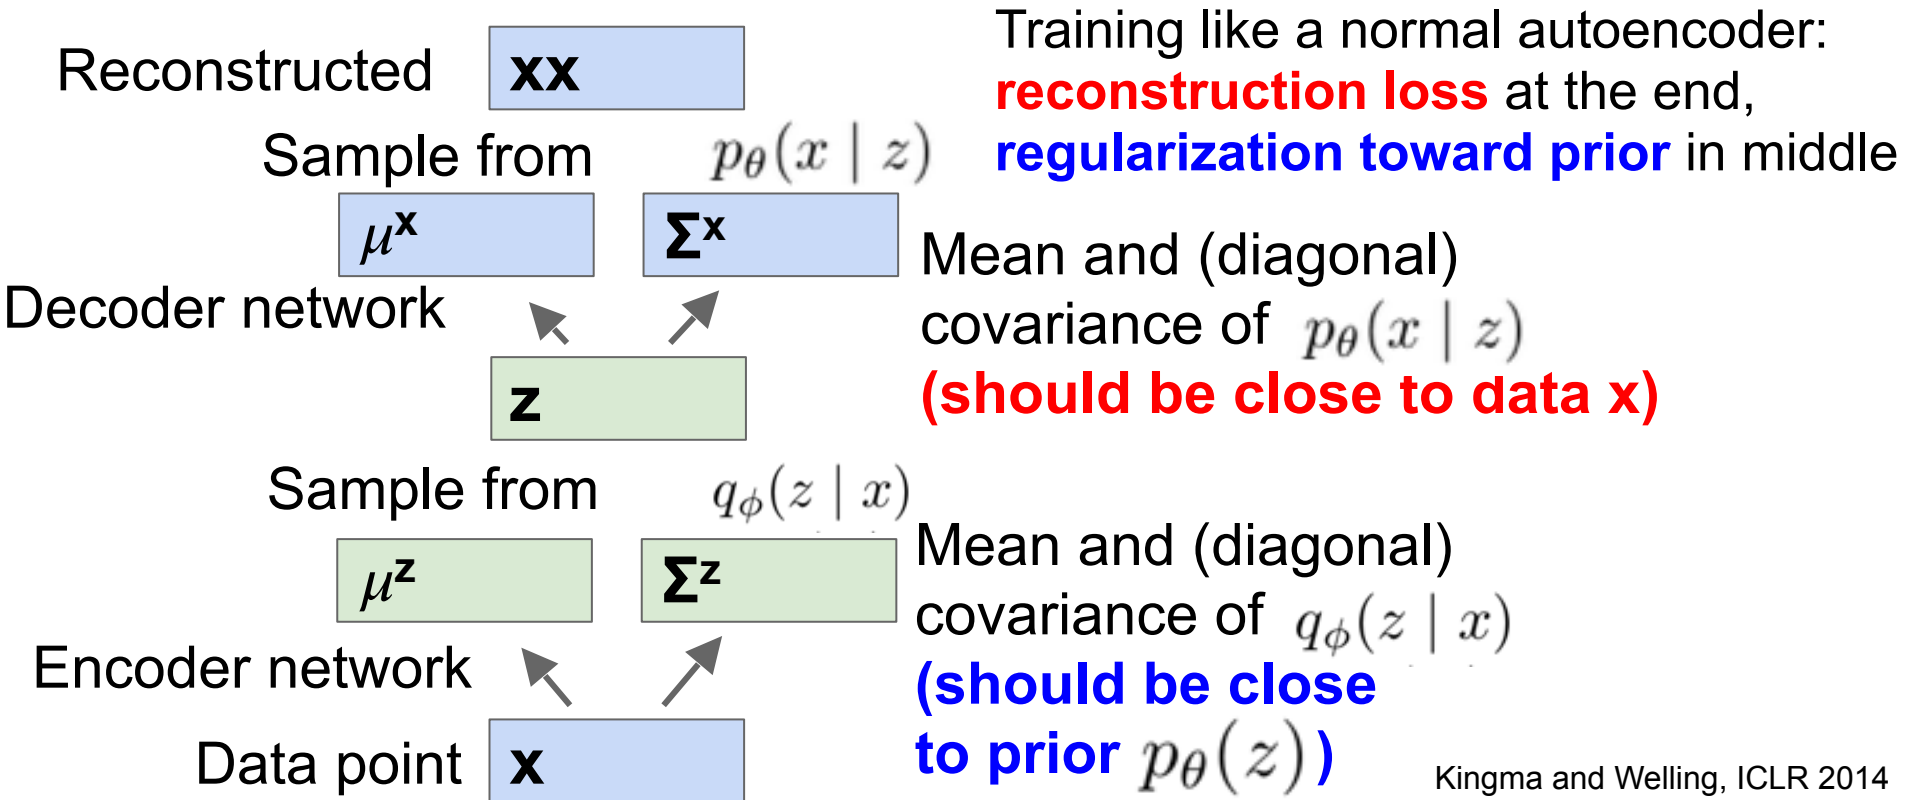
until convergence

Backprop
(Torch7 / Theano)

e.g. Adagrad

Scales to very large datasets!

Variational Autoencoder



Autoencoder Overview

- Traditional Autoencoders
 - Try to reconstruct input
 - Used to learn features, initialize supervised model
 - Not used much anymore
- Variational Autoencoders
 - Bayesian meets deep learning
 - Sample from model to generate images

Generative Adversarial Networks

Generative Adversarial Nets

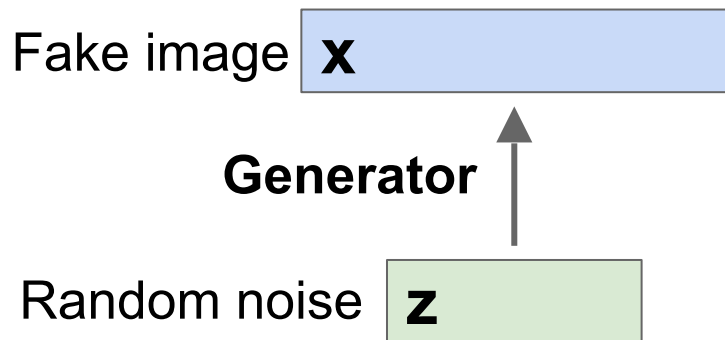
Can we generate images with less math?

Random noise

z

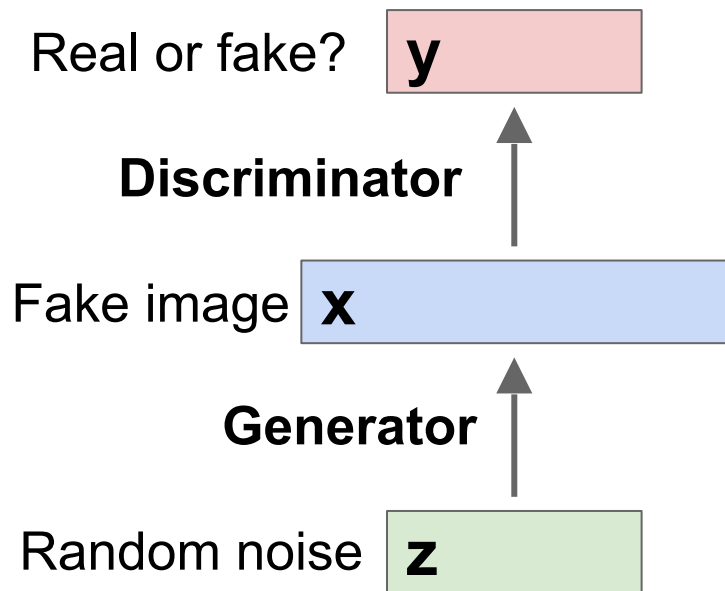
Generative Adversarial Nets

Can we generate images with less math?



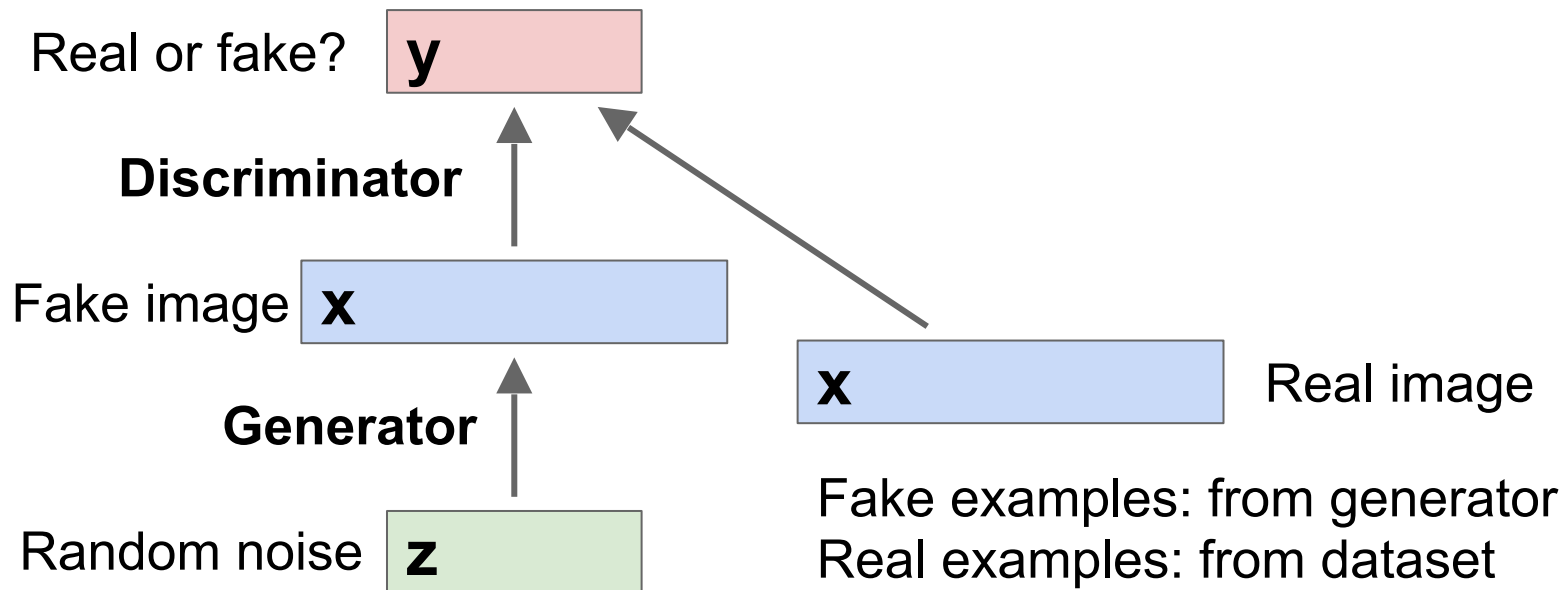
Generative Adversarial Nets

Can we generate images with less math?



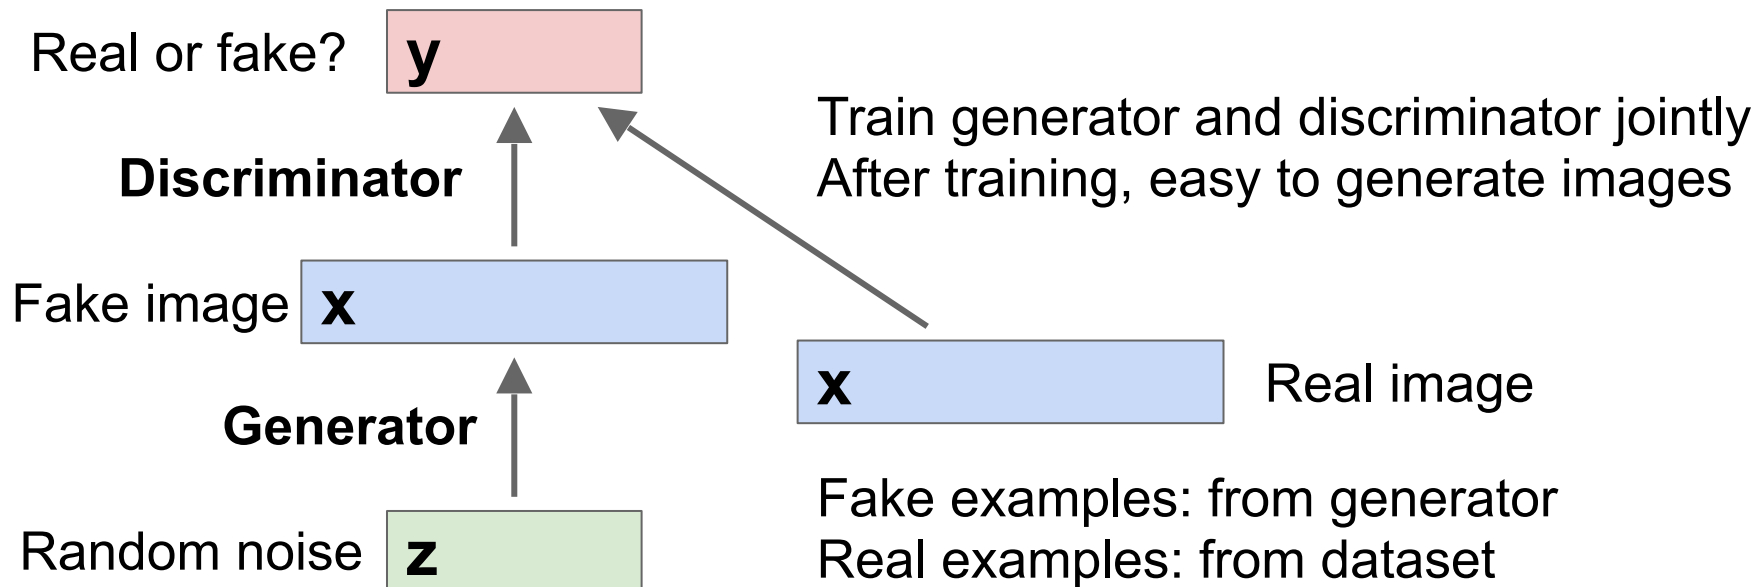
Generative Adversarial Nets

Can we generate images with less math?

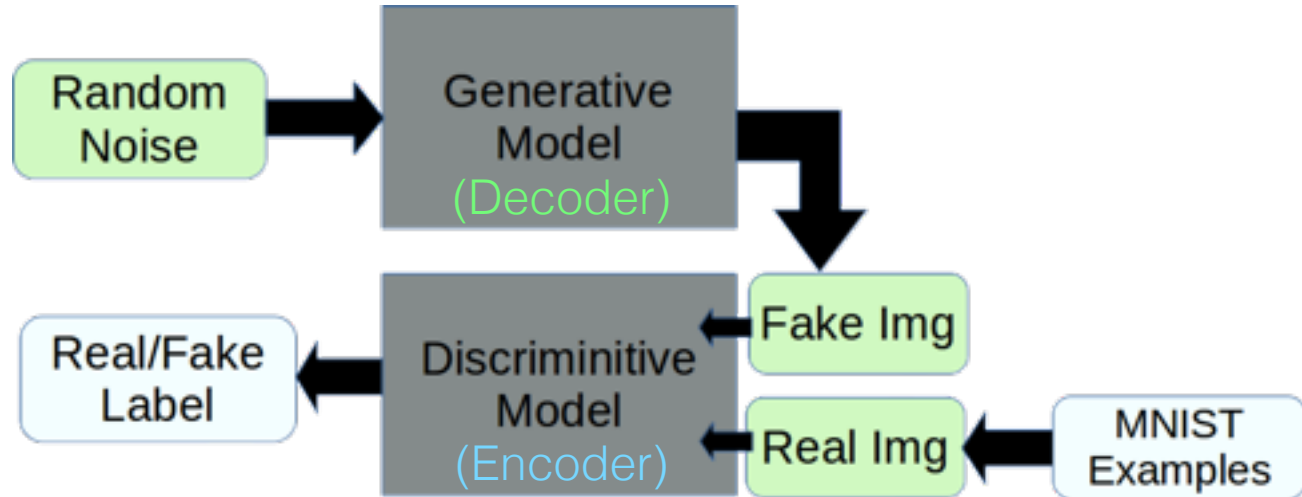


Generative Adversarial Nets

Can we generate images with less math?




Generative Adversarial Nets



Generative Network

Generated
Image




Radford et al, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks", ICLR 2016

Discriminative Network

Real Training
Image

Classified
Label Vector



This is just a CNN!

Generative Adversarial Nets: Simplifying

Samples
from the
model look
amazing!



Radford et al,
ICLR 2016

* Original slides borrowed from Andrej Karpathy
and Li Fei-Fei, Stanford cs231n

Generative Adversarial Nets: Simplifying

Interpolating
between
random
points in
latent space

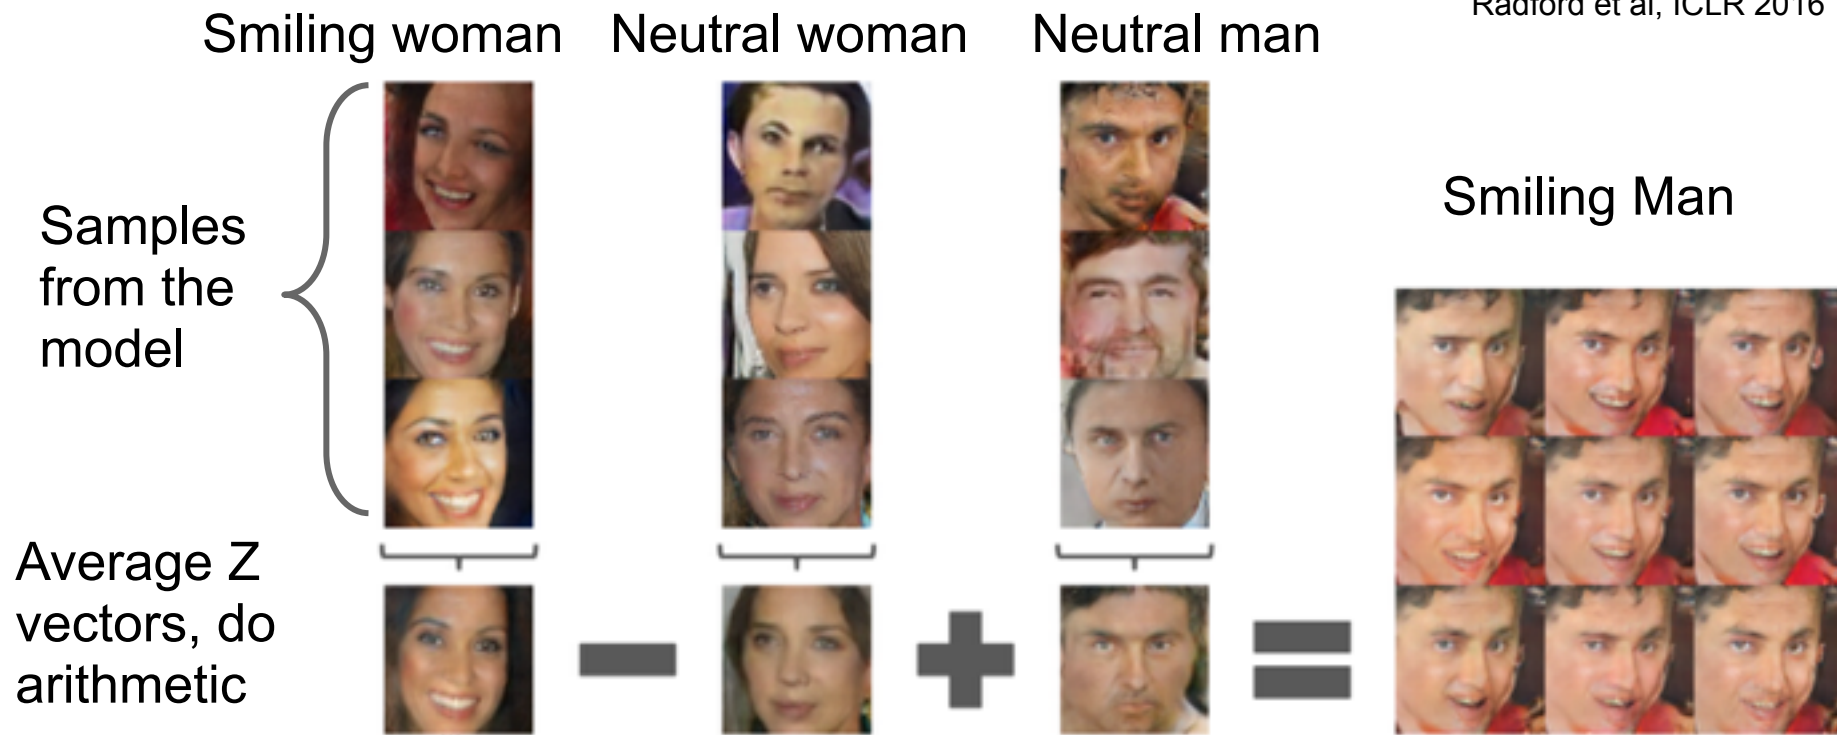


Radford et al,
ICLR 2016

* Original slides borrowed from Andrej Karpathy
and Li Fei-Fei, Stanford cs231n

Generative Adversarial Nets: Vector Math

Radford et al, ICLR 2016



Generative Adversarial Nets: Vector Math

Glasses man

No glasses man

No glasses woman



-

+

=

Woman with glasses



Radford et al,
ICLR 2016

* Original slides borrowed from Andrej Karpathy
and Li Fei-Fei, Stanford cs231n

Learning what to Ignore

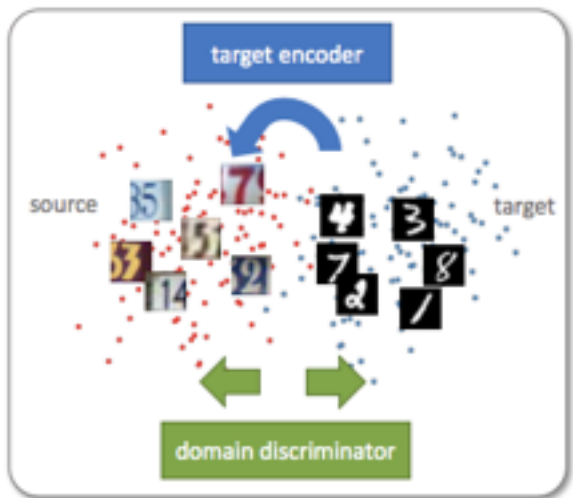


Figure 1: We propose an improved unsupervised domain adaptation method that combines adversarial learning with discriminative feature learning. Specifically, we learn a discriminative mapping of target images to the source feature space (target encoder) by fooling a domain discriminator that tries to distinguish the encoded target images from source examples.



Tzeng et al, “Adversarial Discriminative Domain Adaptation”, arXiv 2017.

Interaction




Figure 9. Interactive image generation and editing. The user can incrementally modify the sketch to change the eyes, hair, and head decorations.

Sangkloy et al, “Scribbler: Controlling Deep Image Synthesis with Sketch and Color”, Siggraph 2017.

Deep Learning and Generalization

(super short) primer on generalization

- given samples $z_1, z_2, \dots, z_n \sim D$, optimize the cost/loss over hypotheses $h \in \mathcal{H}$:

$$\mathbb{E}\{f(z, h)\}$$

- want $h_*(z_1, \dots, z_n)$ such that $\mathbb{E}\{f(z, h_*)\} \leq \inf_h \mathbb{E}\{f(z, h)\} + \epsilon(n)$

Central finding of Zhang et al (2017):

deep neural nets are able to fit random labels and data

So how are Deep Nets
achieving good generalization?


datasets and models

- CIFAR10 dataset:
60000 images (50000 train, 10000 validation), 10 categories
- ImageNet dataset:
1,281,167 training images, 50000 validation images, 1000 categories
- alexnet, inception, multilayer perceptrons


randomization tests:

- true labels: original dataset
- partially corrupted labels: randomize p fraction of labels
- random labels: randomize all labels
- shuffled pixels: uniform random permutation applied to all images
- random pixels: different random permutation
- gaussian: pixels replaced by i.i.d. gaussian r.v.


performance on randomized tests



(a) learning curves



(b) convergence slowdown



(c) generalization error growth

explicit regularization does not help much

Table 1: The training and test accuracy (in percentage) of various models on the CIFAR10 dataset. Performance with and without data augmentation and weight decay are compared. The results of fitting random labels are also included.

model	# params	random crop	weight decay	train accuracy	test accuracy
Inception	1,649,402	yes	yes	100.0	89.05
		yes	no	100.0	89.31
		no	yes	100.0	86.03
		no	no	100.0	85.75
(fitting random labels)		no	no	100.0	9.78
Inception w/o BatchNorm	1,649,402	no	yes	100.0	83.00
		no	no	100.0	82.00
		no	no	100.0	10.12
Alexnet	1,387,786	yes	yes	99.90	81.22
		yes	no	99.82	79.66
		no	yes	100.0	77.36
		no	no	100.0	76.07
(fitting random labels)		no	no	99.82	9.86
MLP 3x512	1,735,178	no	yes	100.0	53.35
		no	no	100.0	52.39
		no	no	100.0	10.48
MLP 1x512	1,209,866	no	yes	99.80	50.39
		no	no	100.0	50.51
		no	no	99.34	10.61